

1 Introduction

The objective of this homework assignment was to understand the purpose of skip blocks and its usage in object detection and localization. With skip connections, a network includes shortcut pathways so that the loss calculated at the output can be felt more strongly in the earlier layers of the network during back-propagation. Layers of convolutional networks and a final linear layer is used to obtain the classification label and predicted bounding box. Two loss functions are necessary to measure the regression loss (i.e. MSE and Complete Box IoU Loss) for the bounding box and cross-entropy loss for the classification label. IoU stands for Intersection over Union and it takes the ratio of the intersection of the two bounding boxes to their union.

2 Methodology

2.1 COCO Downloader

The script below was used to load all the images and save them into its respective categories and store its class label and the provided bounding box into a dataframe. The bounding box saved into the dataframe is that of the one surrounding the dominant object and has an area greater than 200×200 pixels.

```
1 from pycocotools.coco import COCO
2 from PIL import Image
3 import numpy as np
4 import os
5 import pandas as pd
6 import torchvision.transforms as tv
7 from torchvision.io import read_image
8 import argparse
9
10 # Global Variables
11 train_directory = r"/Users/nikitaravi/Documents/Academics/ECE 60146/HW5/
    train2014"
12 val_directory = r"/Users/nikitaravi/Documents/Academics/ECE 60146/HW5/
    val2014"
13 train_annotations = r"/Users/nikitaravi/Documents/Academics/ECE 60146/HW5/
    annotations/instances_train2014.json"
14 val_annotations = r"/Users/nikitaravi/Documents/Academics/ECE 60146/HW5/
    annotations/instances_val2014.json"
15 coco_train = COCO(train_annotations)
16 coco_val = COCO(val_annotations)
17 categories = ["pizza", "bus", "cat"]
18 columns = ["id", "category", "path_to_image", "x1", "y1", "x2", "y2"]
```

```

19 new_image_size = 256
20 resize_image = tvf.Compose([tvf.Resize((new_image_size,new_image_size))])
21
22
23 def parser():
24     parser = argparse.ArgumentParser(description="Load data into dataframe
25 ")
26     parser.add_argument("--train", action="store_true", help="Download
27 training data or testing data")
28     args = parser.parse_args()
29     return args
30
31 def create_dataframe(train=True):
32     saved_count = 0
33     df = pd.DataFrame(columns=columns)
34     ids, class_list, paths_to_image, bbox_x1s, bbox_y1s, bbox_x2s,
35     bbox_y2s = [], [], [], [], [], [], []
36     directory = train_directory if train else val_directory
37     new_dirname = f"./train" if train else f"./val"
38     coco = coco_train if train else coco_val
39     os.makedirs(new_dirname, exist_ok=True)
40     for class_idx, category in enumerate(categories):
41         category_id = coco.getCatIds(catNms=category)
42         image_ids = coco.getImgIds(catIds=category_id)
43         for img_id in image_ids:
44             current_image = coco.loadImgs(ids=img_id)
45             current_id = current_image[0]["id"]
46             annotation_id = coco.getAnnIds(imgIds=img_id, catIds=
47 category_id, iscrowd=False)
48             annotation = coco.loadAnns(ids=annotation_id)
49
50             check_bbox = False
51             if(annotation[0]["area"] > 200 * 200):
52                 saved_count += 1
53                 check_bbox = True
54             if(check_bbox):
55                 path_to_image = current_image[0]["file_name"]
56                 original_width, original_height = current_image[0]["width"
57 ], current_image[0]["height"]
58
59                 # Resize Image
60                 original_image = resize_image(read_image(os.path.join(
61 directory, path_to_image)))
62                 original_image = original_image.repeat(3,1,1) if
63 original_image.size()[0] == 1 else original_image # Check if image has
64 three channels
65                 original_image = tvf.functional.to_pil_image(
66 original_image).convert("RGB") # convert to RGB image
67
68                 # Resize bbox
69                 annotation_width, annotation_height = annotation[0]["bbox"
70 ][2], annotation[0]["bbox"][3]
71                 annotation_left_x, annotation_left_y = annotation[0]["bbox
72 "][0], annotation[0]["bbox"][1]

```

```

62         x_scale = new_image_size / original_width
63         y_scale = new_image_size / original_height
64
65         new_bbox_width = x_scale * annotation_width
66         new_bbox_height = y_scale * annotation_height
67         new_x1 = x_scale * annotation_left_x
68         new_y1 = y_scale * annotation_left_y
69
70         new_x2 = new_x1 + new_bbox_width
71         new_y2 = new_y1 + new_bbox_height
72
73         image_filename = f"{category}_{current_id}.png"
74         original_image.save(os.path.join(new_dirname,
75 image_filename))
76
77         # Add everything to the list to add to the df
78         ids.append(current_id)
79         class_list.append(category)
80         paths_to_image.append(os.path.join(new_dirname,
81 image_filename))
82         bbox_x1s.append(new_x1)
83         bbox_y1s.append(new_x2)
84         bbox_x2s.append(new_y1)
85         bbox_y2s.append(new_y2)
86
87         df["id"] = ids
88         df["category"] = class_list
89         df["path_to_image"] = paths_to_image
90         df["x1"] = bbox_x1s
91         df["x2"] = bbox_y1s
92         df["y1"] = bbox_x2s
93         df["y2"] = bbox_y2s
94
95         filename = "train_data.csv" if train else "test_data.csv"
96         df.to_csv(filename)
97
98 if __name__ == "__main__":
99     args = parser()
100     train = True if args.train else False
101
102     create_dataframe(train)

```

Listing 1: COCO Downloader

2.1.1 Inputs

The following image is a collection of the input images with its corresponding bounding box provided by the COCO dataset

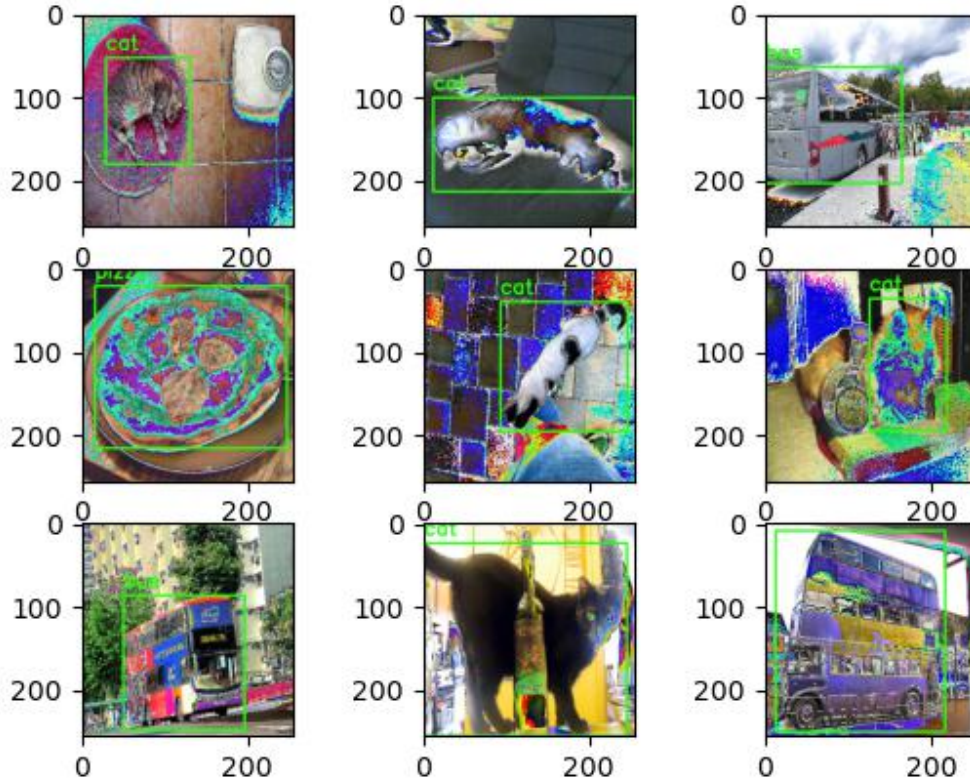


Figure 1: Inputs

2.2 Training

The methodology involved to train and test the model against the training and validation dataset respectively is as follows:

- The loss function used for classifying the images is the cross-entropy loss
- The loss function used for predicting the bounding box is the mean squared error loss or the complete box IoU loss
- The optimizer used is the Adam optimizer

- The loss graph from the cross-entropy loss function is retained to compute the bounding box regression loss with respect to that label
- Every 50 iterations, the model is saved and the average running loss is printed out

2.2.1 Source Code

```

1 # Import Libraries
2 import numpy as np
3 import torch
4 import torchvision.transforms as tv
5 import torch.utils.data
6 import torch.nn as nn
7 import torch.nn.functional as F
8 import matplotlib.pyplot as plt
9 from PIL import Image
10 import os
11 import seaborn as sns
12 from torchvision.ops import complete_box_iou_loss
13 import pandas as pd
14 import cv2
15 import argparse
16
17 import warnings
18 warnings.simplefilter(action='ignore', category=FutureWarning)
19
20 # GLOBAL VARIABLES
21 device = 'cuda' if torch.cuda.is_available() else 'cpu'
22 device = torch.device(device)
23 path_to_model = r"/content/drive/MyDrive/ECE 60146/HW5/model/"
24
25 ##### CREATING DATASETS
26 #####
27 class GenerateDataset(torch.utils.data.Dataset):
28     def __init__(self, df, transform=None):
29         super().__init__()
30         self.df = df
31         self.transform = transform
32
33     def __return_integer_encoding(self, category):
34         categories = {"pizza": 0, "bus": 1, "cat": 2}
35         return categories[category]
36
37     def __len__(self):
38         return len(self.df)
39
40     def __getitem__(self, idx):
41         image_info = self.df.iloc[idx]
42         path_to_image = os.path.join(r"/content/drive/MyDrive/ECE 60146/HW5/", image_info["path_to_image"])
43         image = Image.open(path_to_image)
44         image = self.transform(image) if self.transform else image

```

```

44         label = int(self.__return_integer_encoding(image_info["category"]))
45     )
46     bbox = [image_info["x1"], image_info["y1"], image_info["x2"],
47             image_info["y2"]]
48     # bbox_with_pixel_coords = [bbox[0], bbox[1], bbox[0] + bbox[2],
49     #                             bbox[1] + bbox[3]]
50     bbox_tensor = torch.tensor(bbox, dtype=torch.float)
51
52     return image, label, bbox_tensor
53
54 def get_training_dataloader():
55     # Constants
56     transform = tvn.Compose([tvn.ToTensor(), tvn.Normalize((0.5, 0.5, 0.5),
57     (0.5, 0.5, 0.5))])
58
59     # Create PyTorch Datasets and Dataloader
60     df = pd.read_csv(r"/content/drive/MyDrive/ECE 60146/HW5/train_data.csv")
61     train_dataset = GenerateDataset(df, transform)
62     #print(train_dataset[0][0].shape, train_dataset[0][1], train_dataset
63     #0[2]); quit()
64     trainloader = torch.utils.data.DataLoader(train_dataset, batch_size=2,
65     num_workers=2, shuffle=True)
66     return trainloader
67
68 def get_testing_dataloader():
69     # Constants
70     transform = tvn.Compose([tvn.ToTensor(), tvn.Normalize((0.5, 0.5, 0.5),
71     (0.5, 0.5, 0.5))])
72
73     # Create PyTorch Datasets and Dataloader
74     df = pd.read_csv(r"/content/drive/MyDrive/ECE 60146/HW5/test_data.csv")
75     test_dataset = GenerateDataset(df, transform)
76     testloader = torch.utils.data.DataLoader(test_dataset, batch_size=2,
77     num_workers=2, shuffle=True)
78     return testloader
79
80 ##### NETWORK #####
81 class ResnetBlock(nn.Module):
82     # Inspired by Professor Kak's SkipBlock class
83     def __init__(self, in_ch, out_ch, downsample=False, skip_connections=
84     True):
85         super(ResnetBlock, self).__init__()
86         self.in_ch = in_ch
87         self.out_ch = out_ch
88         self.downsample = downsample
89         self.skip_connections = skip_connections
90
91         self.conv1 = nn.Conv2d(in_ch, out_ch, kernel_size=3, stride=1,
92         padding=1)
93         self.conv2 = nn.Conv2d(in_ch, out_ch, kernel_size=3, stride=1,
94         padding=1)

```

```

85         self.bn1 = nn.BatchNorm2d(out_ch)
86         self.bn2 = nn.BatchNorm2d(out_ch)
87
88         if(self.downsample):
89             self.downsampler = nn.Conv2d(in_ch, out_ch, kernel_size=1,
90 stride=2)
91
92     def forward(self, x):
93         identity = x
94         out = self.conv1(x)
95         out = self.bn1(out)
96         out = torch.nn.functional.relu(out)
97
98         if(self.in_ch == self.out_ch):
99             out = self.conv2(out)
100             out = self.bn2(out)
101             out = torch.nn.functional.relu(out)
102         if(self.downsample):
103             out = self.downsampler(out)
104             identity = self.downsampler(identity)
105         if(self.skip_connections):
106             if(self.in_ch == self.out_ch):
107                 out = out + identity
108             else:
109                 out = torch.cat((out[:, :, self.in_ch :, :] + identity, out[:,
self.in_ch :, :, :] + identity), dim=1)
110         return out
111
112 class HW5Net(nn.Module):
113     # Inspired by Professor Kak's Loadnet2
114     def __init__(self, skip_connections=True, depth=16):
115         super(HW5Net, self).__init__()
116         self.skip_connections = skip_connections
117         self.depth = depth // 2
118         self.conv = nn.Conv2d(in_channels=3, out_channels=64, kernel_size
=3, padding=1)
119         self.pool = nn.MaxPool2d(kernel_size=2, stride=2)
120
121         # Classification
122         self.bn1 = nn.BatchNorm2d(num_features=64)
123         self.bn2 = nn.BatchNorm2d(num_features=128)
124         self.skip64_arr = nn.ModuleList()
125         for idx in range(self.depth):
126             self.skip64_arr.append(ResnetBlock(in_ch=64, out_ch=64,
skip_connections=self.skip_connections))
127             self.skip64ds = ResnetBlock(in_ch=64, out_ch=64, downsample=True,
skip_connections=self.skip_connections)
128             self.skip64to128 = ResnetBlock(in_ch=64, out_ch=128,
skip_connections=self.skip_connections)
129             self.skip128_arr = nn.ModuleList()
130             for idx in range(self.depth):
131                 self.skip128_arr.append(ResnetBlock(in_ch=128, out_ch=128,
skip_connections=self.skip_connections))

```

```

132     self.skip128ds = ResnetBlock(in_ch=128, out_ch=128, downsample=
133     True, skip_connections=self.skip_connections)
134
135     self.fc1 = nn.Linear(in_features=32*32*128, out_features=3)
136     # self.fc2 = nn.Linear(in_features=1000, out_features=3) # Outputs
137     probability of three classes
138
139     # Regression
140     self.conv_seqn = nn.Sequential(nn.Conv2d(in_channels=64,
141     out_channels=64, kernel_size=3, padding=1),
142     nn.BatchNorm2d(num_features=64),
143     nn.ReLU(inplace=True),
144     nn.Conv2d(in_channels=64,
145     out_channels=64, kernel_size=3, padding=1),
146     nn.ReLU(inplace=True))
147
148     self.fc_seqn = nn.Sequential(nn.Linear(in_features=128*128*64,
149     out_features=4)) # Outputs [x, y, x+w, y+h]
150
151     def forward(self, x):
152         x = self.pool(nn.functional.relu(self.conv(x)))
153
154         # Classification
155         cls = x.clone()
156         for idx, skip64 in enumerate(self.skip64_arr[:self.depth//4]):
157             cls = skip64(cls)
158         cls = self.skip64ds(cls)
159         for idx, skip64 in enumerate(self.skip64_arr[self.depth//4:]):
160             cls = skip64(cls)
161         cls = self.bn1(cls)
162         cls = self.skip64to128(cls)
163         for idx, skip128 in enumerate(self.skip128_arr[:self.depth//4]):
164             cls = skip128(cls)
165         cls = self.bn2(cls)
166         cls = self.skip128ds(cls)
167         for idx, skip128 in enumerate(self.skip128_arr[self.depth//4:]):
168             cls = skip128(cls)
169         # print(cls.shape); quit()
170         cls = cls.view(-1, 32 * 32 * 128)
171         cls = self.fc1(cls)
172         #cls = self.fc2(cls)
173
174         # Regression for Bbox
175         bbox = self.conv_seqn(x)
176         # print(bbox.shape); quit()
177         bbox = bbox.view(x.size(0), -1)
178         bbox = self.fc_seqn(bbox)
179
180         return cls, bbox
181
182     ##### TRAINING #####
183     def train(model, trainloader, path_to_model=path_to_model, mse=True, lr=1e
184     -4, betas=(0.9, 0.99), epochs=7):
185         # Train function inspired by Professor Kak's

```



```

run_code_for_training_with_CrossEntropy_and_MSE_Losses function
180 print("Training Started")
181 model = model.to(device)
182 num_layers = len(list(model.parameters()))
183 assert num_layers >= 50, f"number of layers greater than or equal to
50 expected, got: {num_layers}"
184 print(f"Number of layers: {num_layers}")
185
186 model_name = "model_mse.pth" if mse else "model_ciou.pth"
187 path_to_model = os.path.join(path_to_model, model_name)
188
189 cls_criterion = nn.CrossEntropyLoss()
190 reg_criterion = nn.MSELoss() if mse else complete_box_iou_loss
191 optimizer = torch.optim.Adam(model.parameters(), lr=lr, betas=betas)
192 labeling_loss_tally = []
193 regression_loss_tally = []
194
195 for epoch in range(1, epochs+1):
196     print(f"Epoch: {epoch}")
197     running_loss_labeling = 0.0
198     running_loss_regression = 0.0
199     for batch_idx, (inputs, labels, bbox) in enumerate(trainloader):
200         print(f"Batch idx: {batch_idx}")
201         inputs = inputs.to(device)
202         labels = labels.to(device)
203         bbox = bbox.to(device)
204
205         optimizer.zero_grad()
206         outputs = model(inputs)
207         output_label, output_bbox = outputs[0], outputs[1]
208
209         label_loss = cls_criterion(output_label, labels)
210         label_loss.backward(retain_graph=True) # Preserve this loss
211         while calculating bbox loss
212             if(mse):
213                 bbox_loss = reg_criterion(output_bbox, bbox)
214             else:
215                 bbox_loss = reg_criterion(output_bbox, bbox, reduction="
mean")
216         bbox_loss.backward()
217         optimizer.step()
218
219         running_loss_labeling += label_loss.item()
220         running_loss_regression += bbox_loss.item()
221
222         if(batch_idx % 50 == 49):
223             avg_loss_labeling = running_loss_labeling / float(50)
224             avg_loss_regression = running_loss_regression / float(50)
225             labeling_loss_tally.append(avg_loss_labeling)
226             regression_loss_tally.append(avg_loss_regression)
227             print(f"Epoch: {epoch}/{epochs}, Iteration: {batch_idx +
1}: Labeling Loss: {avg_loss_labeling}, Regression Loss:{
avg_loss_regression}")
228             print(f"-----Saving model-----")

```

```

228         torch.save(model.state_dict(), path_to_model)
229
230         running_loss_labeling = 0.0
231         running_loss_regression = 0.0
232
233     return labeling_loss_tally, regression_loss_tally
234
235 def test(model, testloader, path_to_model=path_to_model, mse=True,
num_classes=3):
236     model_name = "model_mse.pth" if mse else "model_ciou.pth"
237     path_to_model = os.path.join(path_to_model, model_name)
238     model.load_state_dict(torch.load(path_to_model))
239     model = model.to(device)
240     confusion_matrix = np.zeros((num_classes, num_classes))
241     image_info = []
242
243     with torch.no_grad():
244         for inputs, labels, bbox in testloader:
245             inputs = inputs.to(device)
246             labels = labels.to(device)
247             bbox = bbox.to(device)
248
249             outputs = model(inputs)
250             output_label = outputs[0]
251             output_bbox = outputs[1].tolist()
252
253             _, predicted = torch.max(output_label, dim=1)
254             for label, prediction in zip(labels, predicted):
255                 confusion_matrix[label][prediction] += 1
256
257             for img, original_label, predicted_label, original_bbox,
predicted_bbox in zip(inputs, labels, predicted, bbox, output_bbox):
258                 image_info.append({"Image": img, "Original Label":
original_label, "Predicted Label": predicted_label,
259                                 "Original Bbox": original_bbox, "
Predicted Bbox": predicted_bbox})
260
261     accuracy = np.trace(confusion_matrix) / np.sum(confusion_matrix)
262     return confusion_matrix, accuracy, image_info
263
264 def plot_labeling_loss(labeling, epochs):
265     iterations = range(len(labeling))
266     figure = plt.figure(1)
267     plt.plot(iterations, labeling, label="Labeling Loss")
268
269     plt.title(f"Loss per Iteration")
270     plt.xlabel(f"Iterations over {epochs} epochs")
271     plt.ylabel("Loss")
272
273     filename = "label_loss.jpg"
274     plt.savefig(os.path.join("/content/drive/MyDrive/ECE 60146/HW5/Results
", filename))
275
276 def plot_regression_loss(regression, epochs, mse):

```

```

277     iterations = range(len(regression))
278     figure = plt.figure(2)
279     plt.plot(iterations, regression, label="Regression Loss")
280
281     plt.title(f"Loss per Iteration")
282     plt.xlabel(f"Iterations over {epochs} epochs")
283     plt.ylabel("Loss")
284
285     filename = "mse_loss.jpg" if mse else "ciou_loss.jpg"
286     plt.savefig(os.path.join("/content/drive/MyDrive/ECE 60146/HW5/Results", filename))
287
288 def display_confusion_matrix(conf, accuracy, class_list=["pizza", "bus", "cat"], mse=True):
289     figure = plt.figure(3)
290     sns.heatmap(conf, xticklabels=class_list, yticklabels=class_list, annot=True)
291     plt.xlabel(f"True Label \n Accuracy: {accuracy}")
292     plt.ylabel("Predicted Label")
293
294     filename = "conf_mse.jpg" if mse else "conf_ciou.jpg"
295     plt.savefig(os.path.join("/content/drive/MyDrive/ECE 60146/HW5/Results", filename))
296
297 def draw_rectangles(image, bbox, predicted_bbox, label, predicted_label):
298     # categories = {"pizza": 0, "bus": 1, "cat": 2}
299     inverse_categories = {0: "pizza", 1: "bus", 2: "cat"}
300
301     label = inverse_categories[int(label.item())]
302     predicted_label = inverse_categories[int(predicted_label.item())]
303
304     image = np.asarray(tvt.ToPILImage()(image))
305     image = cv2.rectangle(image, (int(bbox[0]), int(bbox[1])), (int(bbox[2]), int(bbox[3])), color=(36, 255, 12), thickness=2)
306     image = cv2.putText(image, label, (int(bbox[0]), int(bbox[1] - 10)), fontFace=cv2.FONT_HERSHEY_SIMPLEX, fontScale=0.8, color=(36, 255, 12), thickness=2)
307
308     image = cv2.rectangle(image, (int(predicted_bbox[0]), int(predicted_bbox[1])), (int(predicted_bbox[2]), int(predicted_bbox[3])), color=(255, 45, 12), thickness=2)
309     image = cv2.putText(image, predicted_label, (int(predicted_bbox[0]), int(predicted_bbox[1] - 10)), fontFace=cv2.FONT_HERSHEY_SIMPLEX, fontScale=0.8, color=(255, 45, 12), thickness=2)
310     return image
311
312 def display_bbox(image_info, mse=True):
313     fig, ax = plt.subplots(3, 3)
314     row, col = 0, 0
315     for idx, details in enumerate(image_info):
316         image = details["Image"]
317         original_label = details["Original Label"]
318         predicted_label = details["Predicted Label"]
319         original_bbox = details["Original Bbox"]

```

```

320     predicted_bbox = details["Predicted Bbox"]
321
322     image = draw_rectangles(image, original_bbox, predicted_bbox,
original_label, predicted_label)
323     ax[row, col].imshow(image)
324     row = row+1 if not ((idx + 1) % 3) else row
325     col = 0 if not((idx + 1) % 3) else col+1
326
327     if(row == 3 and col == 0):
328         break
329
330     filename = "bbox_mse.jpg" if mse else "bbox_ciou.jpg"
331     plt.savefig(os.path.join("/content/drive/MyDrive/ECE 60146/HW5/Results
", filename))
332
333 def display_input_bbox(dataloader):
334     inverse_categories = {0: "pizza", 1: "bus", 2: "cat"}
335     fig, ax = plt.subplots(3, 3)
336     row, col = 0, 0
337     for batch_idx, (images, labels, bboxes) in enumerate(dataloader):
338         for idx in range(len(labels)):
339             if(col == 3):
340                 row += 1
341                 col = 0
342             if(row == 3):
343                 break
344             image, label, bbox = images[idx], inverse_categories[int(
labels[idx])], bboxes[idx]
345             image = np.asarray(tvt.ToPILImage()(image))
346             image = cv2.rectangle(image, (int(bbox[0]), int(bbox[1])), (
int(bbox[2]), int(bbox[3])), color=(36, 255, 12), thickness=2)
347             image = cv2.putText(image, label, (int(bbox[0]), int(bbox[1] -
10)), fontFace=cv2.FONT_HERSHEY_SIMPLEX, fontScale=0.8, color=(36,
255, 12), thickness=2)
348             ax[row, col].imshow(image)
349             col += 1
350             if(row == 3):
351                 break
352
353     plt.savefig(os.path.join("/content/drive/MyDrive/ECE 60146/HW5/Results
", "inputs.jpg"))
354
355 def parser():
356     parser = argparse.ArgumentParser(description="Object Detection and
Localization")
357     parser.add_argument("--epochs", type=int, help="Number of epochs",
default=8)
358     parser.add_argument("--mse", action="store_true", help="Choosing MSE
Loss")
359     args = parser.parse_args()
360     return args
361
362
363 if __name__ == "__main__":

```

```

364     args = parser()
365     epochs = args.epochs
366     regression = True if args.mse else False
367
368     trainloader = get_training_dataloader()
369     testloader = get_testing_dataloader()
370     display_input_bbox(trainloader)
371
372     model = HW5Net()
373     label_loss, regression_loss = train(model, trainloader, epochs=epochs,
374                                       mse=regression)
375     plot_labeling_loss(label_loss, epochs=epochs)
376     plot_regression_loss(regression_loss, epochs=epochs, mse=regression)
377
378     confusion_matrix, accuracy, image_info = test(model, testloader, mse=
379     regression)
380     display_confusion_matrix(confusion_matrix, accuracy, mse=regression)
381     display_bbox(image_info, mse=regression)

```

Listing 2: Training

3 Results

Number of layers: 172

Figure 2: Number of Layers in the Network




Figure 3: Cross Entropy Loss for Classification




Figure 4: MSE Loss for Bounding Box

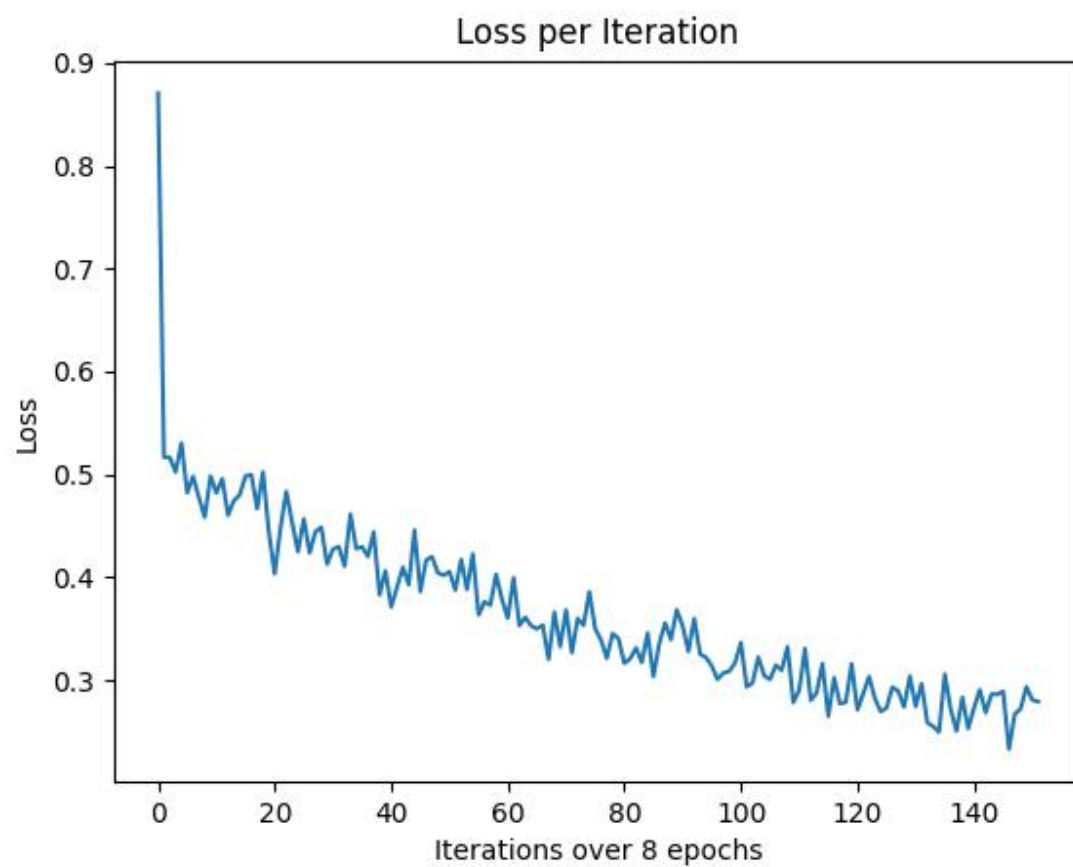


Figure 5: Complete IOU Loss




Figure 6: Confusion Matrix when using MSE Loss

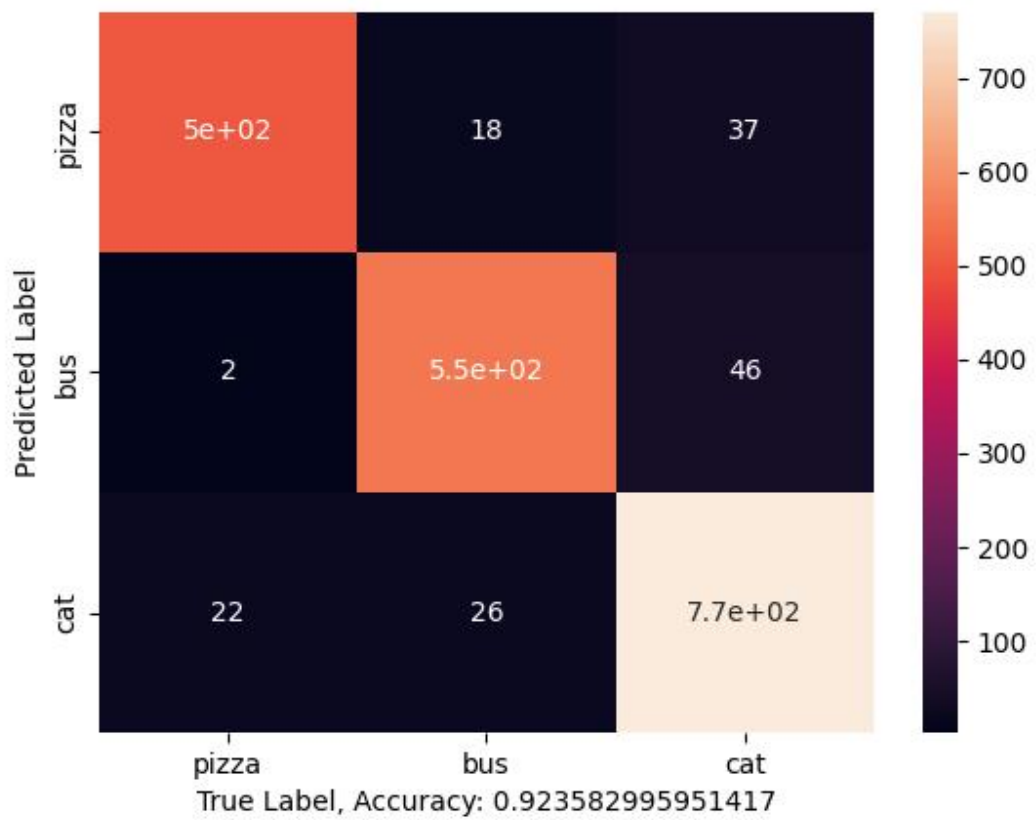


Figure 7: Confusion Matrix when using Complete IOU Loss

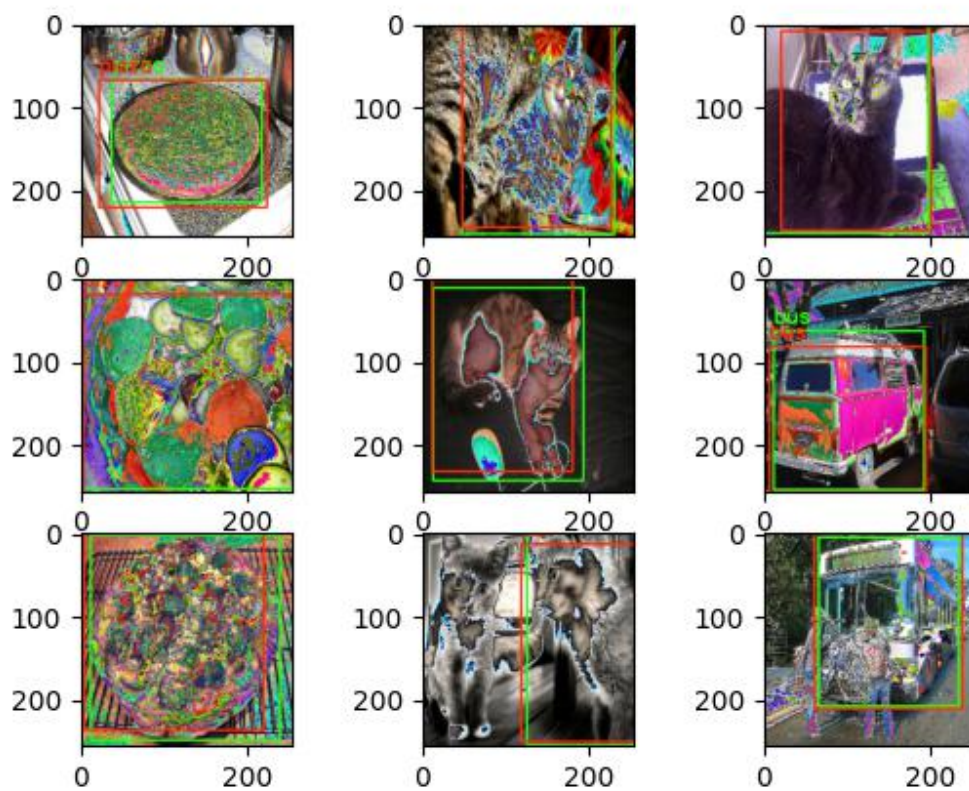


Figure 8: Bounding Box Predictions using MSE Loss

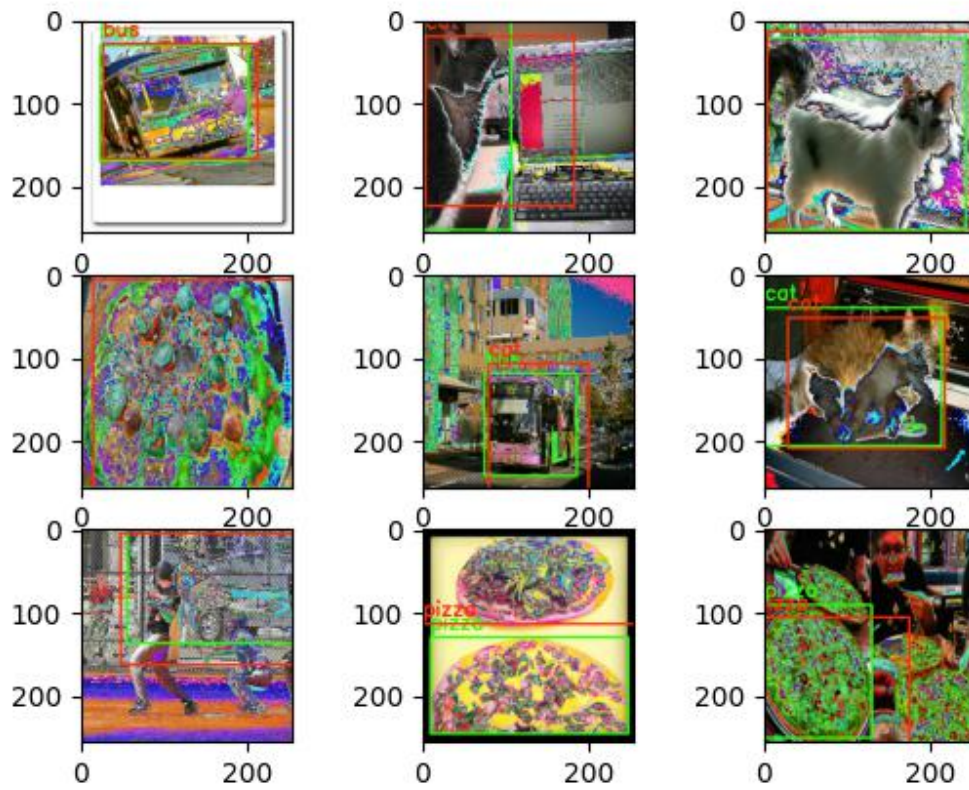


Figure 9: Bounding Box Predictions using Complete IOU Loss

4 Evaluation

This model consists of 172 layers. Using the MSE loss, the accuracy obtained was 89.7%, while using the Complete IoU loss gave an accuracy of 92.4%. In this particular situation, the Complete IoU loss function slightly outperformed the MSE loss function. In both cases, the bounding boxes are slightly off from the actual bounding box. Some enhancements that can be made to this model are:

- Adding more convolutional layers
- Revising the skip block network to something similar to the one found in Resnet
- Changing the learning rate
- Increasing the number of epochs